

Plant Disease Detection and Classification using Deep Learning

1. Abstract

This study assesses the effectiveness of three CNN models—GoogleNet, MobileNetV2, and EfficientNet—in detecting and classifying plant diseases using images from different plant species. The models were tested on three datasets of varying sizes and classes, employing both training-from-scratch and transfer learning methods. The analysis focuses on the influence of data quality, preprocessing techniques, and hyperparameter adjustments on model performance, using metrics such as Accuracy, Precision, Recall, and F1 Score. EfficientNet showed competitive results across all datasets, with accuracies of 64.96% on Cassava, 78.38% on Crop Disease, and 98.61% on Plant Village. GoogleNet achieved the highest accuracy on the Crop Disease dataset at 78.92%, while MobileNetV2 excelled on the Cassava and Plant Village datasets with accuracies of 66.55% and 99.46%, respectively. Transfer learning improved feature detection in images benefitting from pre-trained weights for learning complex embedded features. Additionally, tuning of hyperparameters (batch size and learning rate) enhanced EfficientNet's performance by 5% on the Cassava dataset.

2. Introduction & Problem Statement

2.1. Problem Statement and Challenges

Agriculture is a major sector that impacts society and the environment while defining the economy of a nation, this necessitates precise plant disease identification for timely intervention. Plant disease recognition poses significant challenges due to the diverse range of symptoms, complicating manual categorization by the human eye. Plant diseases can show a range of symptoms, including color changes, spots, wilting, and deformities. These symptoms are affected by factors like the type of pathogen and the environment. These diseases often escape manual identification as they may not exhibit symptoms in the early stages, necessitating the development of advanced techniques for early detection. Monitoring for disease outbreaks over large agricultural areas is hard with manual methods because there's just so much land to cover. Therefore, it is important to develop automated solutions that can aid in identifying crop diseases to timely manage disease infestation.

Convolutional Neural Networks have proven to be effective for computer vision tasks like classification, detection, recognition, etc. due to their ability to perform embedded feature extraction. However, implementing CNNs can be quite complex, as there are various challenges involved. Overfitting is a significant concern with Deep Convolutional Neural Networks (Deep CNNs), given their large number of parameters, it can hinder generalization and lead to poor performance. Additionally, selecting the appropriate archi-

ture for a given problem, considering factors like image type, dataset size, and specific task requirements, is crucial for achieving optimal results. Moreover, image quality issues such as variations in lighting, growth stages, and environmental factors can introduce further obstacles, often leading to incorrect classification. Addressing these challenges requires a well-analyzed thorough approach to model design and training, emphasizing the importance of finding common ground to navigate through these complexities effectively.

In existing solutions, these problems have been addressed through various mechanisms. Data augmentation techniques like rotation, flipping, resizing, and scaling of images, along with background suppression, are commonly employed to diversify and expand the training dataset. These augmentations not only enhance model robustness but also contribute to improved generalization. Moreover, the transfer learning approach, which leverages pre-trained weights, has emerged as a valuable methodology for various computer vision tasks. By fine-tuning these models for specific plant disease detection tasks, transfer learning mitigates the need for extensively annotated datasets while enhancing overall model performance. Through these integrated approaches, existing solutions have enhanced model performance and minimized information leakage risks [13] [14]. However, pre-trained models may not always capture all relevant features for a specific domain such as plant disease detection. Both data augmentation and transfer learning techniques can require significant computational resources, especially when dealing with large datasets and complex models.

Our study involves training nine models from scratch and two via Transfer Learning on three datasets. Using Cross Entropy Loss and SGD Optimizer, models were trained for 20 epochs with a batch size of 32 and a learning rate of 0.001. Evaluation metrics include Accuracy, Precision, Recall, F1-Score, and continuous monitoring of these metrics enabled adjustments for optimal learning. Overall, the performance of the 3 architectures varies across datasets, with EfficientNet showing consistent competitive performance across all, making it a good choice for solving our problem. It is also observed that all the architectures achieve higher performance on the Plant Village dataset as images feature lower background noise. In terms of computational efficiency, MobileNetV2 tends to achieve a minimum time per epoch across all datasets.

2.2. Literature Review and Related Works

In recent years, numerous studies have examined the efficacy of different deep learning models in identifying plant diseases from leaf images. Nandi et al. conducted

000	054
001	055
002	056
003	057
004	058
005	059
006	060
007	061
008	062
009	063
010	064
011	065
012	066
013	067
014	068
015	069
016	070
017	071
018	072
019	073
020	074
021	075
022	076
023	077
024	078
025	079
026	080
027	081
028	082
029	083
030	084
031	085
032	086
033	087
034	088
035	089
036	090
037	091
038	092
039	093
040	094
041	095
042	096
043	097
044	098
045	099
046	100
047	101
048	102
049	103
050	104
051	105
052	106
053	107

108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
a study [13] to develop a device-friendly diagnostic system for guava plant disease detection. They compared the performance of VGG-16, GoogleNet, ResNet-18, MobileNet-v2, and EfficientNet across a dataset comprising images of guava plants. For deployment, they employed Float16 and Dynamic Range Quantization to reduce model size. Their evaluation revealed that EfficientNet attained the highest overall accuracy at 99%. However, quantized GoogleNet, with its size of 0.143 MB and accuracy of 97%, emerged as the optimal candidate for deployment, particularly suitable for devices with limited storage capacity. Bouleit et al. conducted a comprehensive survey [14] comprising 19 studies, aimed at elucidating the advantages and utilization of Convolutional Neural Networks (CNNs) in crop disease identification. Their findings underscored the necessity of a diverse and large image dataset for achieving high performance with CNNs. They emphasized the importance of employing transfer learning, fine-tuning, and careful architecture selection in the development of efficient models. Additionally, the survey highlighted the promotion of regularization techniques and data augmentation to bolster model generalization capabilities.

Shelar et. al. explain a sophisticated architecture for their image classification task, specifically employing the VGG-19 architecture, known for its depth and effectiveness in image classification tasks using multiple layers [12]. This depth allows it to capture a wide range of features from input images, which is crucial for plant disease identification. They also highlighted the importance of deep learning in agriculture, offering a scalable solution for both all sizes of operations. Also, in Sharma et al.'s study, the performance of Convolutional Neural Networks (CNNs) was investigated for identifying diseases in rice and potato crops [15]. They employed a model trained from scratch, comprising convolutional layers, batch normalization, max-pooling, dropout, and fully connected layers. Their model achieved high classification accuracies of 99.58% for rice leaf diseases and 97.66% for potato leaf diseases, with resizing being the only pre-processing step applied.

In our methodology, we drew inspiration from various works, including those mentioned above, and opted to utilize GoogleNet, EfficientNet, and MobileNetV2 architectures. To adapt these architectures to our specific problem, we modified the final fully connected layer. Additionally, we employed transfer learning, fine-tuning, and data augmentation techniques in our approach.

3. Methodology

Our study focuses on using three CNN architectures that were carefully chosen for their optimal balance between size and performance. We conducted a series of experiments, training nine models across various scenarios and employing transfer learning on two models to com-

pare performance with training-from-scratch methods. Additionally, we focused on extensive hyper-parameter tuning for the EfficientNet architecture using the Cassava Dataset. Through adjustments in learning rate and batch size, we aimed to identify the optimal settings for enhanced efficiency and accuracy.

3.1. Datasets

For solving the problem discussed, we have selected the *Cassava Dataset*, *Crop Disease dataset*, and *Plant Village Dataset* as shown in Table 1. The Cassava Dataset contains leaf images of the cassava plant crowd sourced from farmers and analysed by experts at the National Crops Resources Research Institute (NaCRRI) and AI lab in Makarere University, Kampala [1]. The Crop Diseases dataset contains leaf images sourced from farms in Ghana with multiple classes of different diseases [2]. The PlantVillage dataset contains healthy and unhealthy leaf images divided into various classes across multiple species and diseases [3].

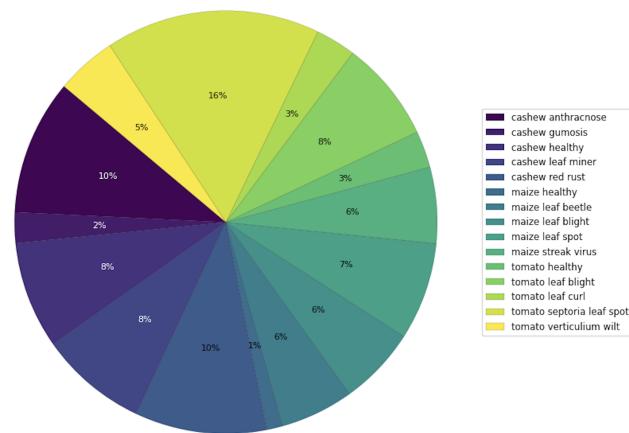


Figure 1. Inter-Class variance of Crop Disease Dataset

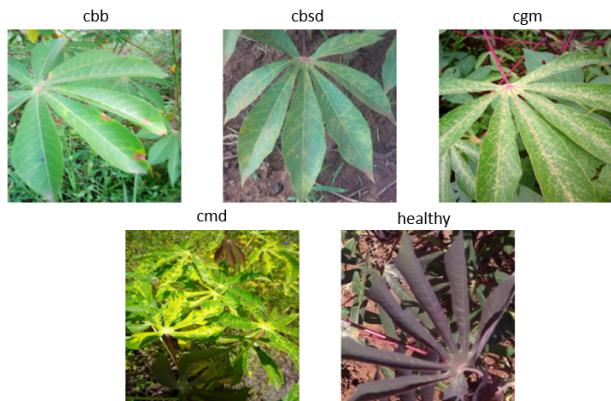


Figure 2. Samples of Cassava Dataset

The inter-class variance of Crop Disease dataset is illus-

162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215

216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
trated in Figure 1 and we see that Plant Village dataset (Figure 7) also has a similar distribution of images among the classes. In contrast, the Cassava dataset (Figure 8) has an uneven distribution with Cassava Mosaic Disease images contributing to almost half of the dataset. The samples of the different images of the classes are illustrated in Figures 2, 10 and 11.

In our project, we have allocated the data as follows: 80% to training, 10% to validation, and 10% to testing, across all datasets. The solution pipeline begins with data pre-processing, given that the datasets exhibit varying geometric and photo-metric properties. Initially, we resized all images to a uniform resolution of 224x224x3, ensuring compatibility with the chosen architectures. Additionally, we applied geometric transformations, such as horizontal and vertical flips, to 10% of the images. Photo-metric transformations were also performed by adjusting the brightness and contrast of the images by a factor of 0.2. Finally, we normalized the images using the ImageNet mean and standard deviation values [5].

	Cassava	Crop Disease	PlantVillage
Images	5,656	24,881	50,091
Size	Var	400x400x3	256x256x3
Color Space	RGB	RGB	RGB
Format	JPG	JPG	JPG
Classes	5	22	37
Trimmed Datasets			
Images	5,656	16,704	35,320
Classes	5	15	30

247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
Table 1. Agricultural Datasets Summary

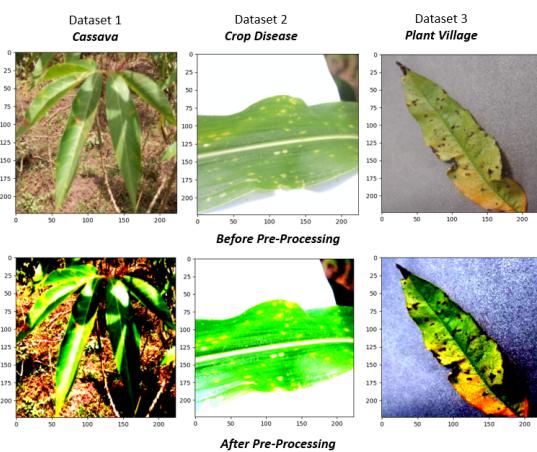


Figure 3. Image samples before and after data augmentation

These pre-processing steps, as seen in Figure 3, including resizing, applying geometric and photo-metric transformations, and normalization according to ImageNet stan-

dards, have significantly enhanced model compatibility and robustness, thereby improving generalization across diverse datasets.

3.2. CNN Models

To address the challenges described above, we have used GoogleNet, MobileNetV2, and EfficientNet B0. Consistent architectural modifications were applied across all CNN models, specifically focusing on their final output layers. This standardization ensures that the output dimensions from the models' last layers correspond to the number of classes present in the dataset. To achieve this, three linear layers were introduced to map the output dimensions sequentially to 512, then 256, and finally to the number of classes. Both BatchNorm and Relu layers were incorporated before and after each linear layer, respectively, to enhance model performance and convergence.

MobileNetV2 has depthwise separable convolution layers due to which the complexity cost and size of the network are reduced, making it useful for devices with low computation power. This model utilizes inverted residuals with linear bottlenecks to maintain representational power and achieves state-of-the-art performance for semantic segmentation along with object detection. We chose MobileNetV2 as it excels in computational efficiency, which is advantageous for resource-constrained environments. MobileNetV2 consists of 53 layers and 3.01 million parameters [22]. It has 0.61 billion FLOPS and took an average of 92s per epoch on all datasets. MobileNetV2 surpasses both MobileNetV1 and ShuffleNet (1.5) in performance, despite having similar model size and computational requirements. When equipped with a width multiplier of 1.4, MobileNetV2 (1.4) outperforms ShuffleNet ($\times 2$) and NASNet while achieving quicker inference times [17].

GoogleNet is a 22-layer (without pooling layers) deep CNN architecture having the inception module as its core building block incorporating multiple filter sizes (1x1, 3x3, 5x5) within a single module. This enables the network to learn features of different scales in the images, improving its ability to capture complex patterns [23]. GoogleNet heavily utilizes 1x1 convolutions for dimensionality reduction, a global average pooling layer for spatial reduction, and auxiliary classifiers to address the vanishing gradient problem to improve overall training stability which makes it suitable for our application. GoogleNet consists of 6.26 million parameters, has 3.01 billion FLOPS, and took an average of 104s per epoch on all datasets. As compared to other models, it allows for parallel processing of information at different spatial scales which helps capture both local and global features effectively and strikes a balance between accuracy and computational efficiency.

EfficientNet handles the trade-off between accuracy, computational efficiency, and model size very well with

GROUP Q: FINAL REPORT

324 the help of compound scaling. It is a concept that focuses
 325 on scaling and balancing three dimensions: width, depth,
 326 and resolution. It consists of MBConv blocks, inspired
 327 by MobileNetV2, which utilizes the squeeze-to-excitation
 328 technique to focus on informative features that also make
 329 it suitable for our use. The baseline network model, B0,
 330 comprises around 18 layers making it relatively lightweight
 331 compared to larger variants in the series [21]. EfficientNet
 332 consists of 4.8 million parameters, has 0.79 billion FLOPS,
 333 and takes an average of 99s per epoch on all datasets. Effi-
 334 cientNet models exhibit either competitive or superior per-
 335 formance on diverse benchmark datasets and computer vi-
 336 sion tasks, making them excellent pre-trained models for
 337 transfer learning, particularly when trained on large-scale
 338 datasets such as ImageNet [19] [20].

339 GoogleNet and MobileNetV2 were chosen for transfer
 340 learning as these models support fine-tuning which enabled
 341 us to tweak and optimize the models' parameters according
 342 to our target task. By leveraging the pre-trained weights, we
 343 could adapt these models to the unique features and char-
 344 acteristics of our dataset resulting in an enhanced perfor-
 345 mance. Overall, the models which are popular and effec-
 346 tive in image classification tasks were chosen while making
 347 sure to address plant disease detection considering compu-
 348 tational efficiency and performance. More information on
 349 time/epoch for all models is given in Table 2.

351 3.3. Optimization Algorithm

352 In our study, the choice of optimization algorithm was
 353 critical to effectively train the CNN models, considering
 354 the challenges posed by the complex image data and the
 355 need for efficient computational performance. We have pri-
 356 marily used Stochastic Gradient Descent (SGD) as our opti-
 357 mization algorithm across all experiments which has helped
 358 maintain a good training and validation accuracy over mul-
 359 tiple epochs.

360 We have chosen SGD for its robustness and effective-
 361 ness in handling large datasets and its ability to converge
 362 even with noisy gradients, which is seen in real-world im-
 363 age data like those in our datasets. While batch gradient de-
 364 scent, which computes the gradient using the entire dataset,
 365 SGD updates the parameters more frequently with gradi-
 366 ents calculated from randomly selected subsets of the data
 367 or mini-batches [24]. We have set the Learning Rate to
 368 0.001, which is a standard value allowing for gradual but
 369 steady learning. While performing hyper-parameter tun-
 370 ing, we have also sampled learning rates uniformly from
 371 the range 0.0001 to 0.001 and chosen the best configura-
 372 tion. The momentum coefficient of 0.9 was also used to
 373 help accelerate the SGD in the relevant direction, thereby
 374 avoiding local minima and faster convergence. A batch size
 375 of 32 was used across the datasets while a study on differ-
 376 ent batch sizes of 8, 16, 32, and 64 was tested for hyper-

377 parameter tuning. Furthermore, we assessed the model's
 378 performance using metrics such as F1-score, Precision, Re-
 379 call, and Accuracy. By leveraging the optimization algo-
 380 rithm, we achieved better accuracy rates, showcasing the
 381 significance of optimization algorithms in enhancing model
 382 performance. We can see the Accuracy and Learning Loss
 383 plots of the different architectures across datasets in Figures
 384 12-33.

385 4. Results

386 4.1. Experimental Setup

387 The experiments were run on the local with Python 3.9
 388 and PyTorch 2.2. We leveraged both an NVIDIA RTX
 389 3080 Ti GPU and Apple M2 Silicon for acceleration. Three
 390 CNN architectures, GoogLeNet, EfficientNet B0, and Mo-
 391 bileNetV2, were used for experimentation. The datasets
 392 were split in the ratio of 80:10:10 for training, validation,
 393 and testing, respectively.

394 In the main study, we have trained 9 models from scratch
 395 using these architectures over 3 datasets. Each model has
 396 been run on 20 Epochs with a batch size of 32, a learn-
 397 ing rate of 0.001, and a momentum of 0.9. We utilized
 398 Stochastic Gradient Descent (SGD) to find optimal weights
 399 and evaluated our models using Cross Entropy as the Loss
 400 Function. Additionally, we have also trained GoogLeNet
 401 and MobileNetV2 architecture using the Transfer Learning
 402 approach on the Crop Diseases dataset with the same exper-
 403 imental setting as earlier.

404 In our ablation study, we investigated the impact of data
 405 augmentation techniques, and hyperparameter tuning on the
 406 performance of the EfficientNet architecture on the Cassava
 407 dataset. We chose the Cassava dataset because its images
 408 contain background noise, which adds complexity to the
 409 classification task. Additionally, the EfficientNet architec-
 410 ture was selected due to its consistent performance in our
 411 main study. We experimented with batch sizes of 8, 16, 32,
 412 and 64, while the learning rate was uniformly sampled from
 413 a range of 0.0001 to 0.001. We utilized Ray 2.10 to track
 414 our experiments with 10 randomly sampled batch size and
 415 learning rate configurations.

416 The performance of models was also evaluated us-
 417 ing accuracy, precision, recall, and F1 Score. Since the
 418 datasets were imbalanced, we employed macro-averaging
 419 techniques while computing precision, recall, and F1 Score.
 420 This approach gives equal weight to each class. Addition-
 421 ally, we also generated Confusion Matrices to gain a deeper
 422 insight into how often a particular class (ground truth) gets
 423 confused with other classes.

424 4.2. Main Results

425 Table 2 summarizes the performance of the chosen ar-
 426 chitectures, trained without pre-trained weights, on un-

GROUP Q: FINAL REPORT

seen images from the selected datasets. We observe that the performance of the three architectures varies across datasets. However, EfficientNet shows consistent competitive performance and generalizes better across all datasets in terms of accuracy, precision, and recall. This is likely due to its architectural design, which utilizes the squeeze-and-excitation technique in MBConv blocks and compound scaling. These features enable EfficientNet to capture complex features effectively.

Dataset	Metrics	GoogleNet	EfficientNet	MobileNetV2
Cassava	Accuracy (%)	59.65	64.96	66.55
	Precision (%)	51.39	51.09	58.14
	Recall (%)	45.20	47.84	44.13
	F1 Score	0.46	0.49	0.42
	Time/Epoch (sec)	45	41	42
Crop Disease	Accuracy (%)	78.92	78.38	74.07
	Precision (%)	77.81	76.46	73.43
	Recall (%)	79.63	76.98	70.33
	F1 Score	0.78	0.76	0.71
	Time/Epoch (sec)	101	97	83
Plant Village	Accuracy (%)	98.13	98.61	99.46
	Precision (%)	98.17	98.80	99.48
	Recall (%)	97.85	98.61	99.20
	F1 Score	0.98	0.98	0.99
	Time/Epoch (sec)	167	161	153

Table 2. Performance of 3 Architectures without Pre-trained weights on different Datasets

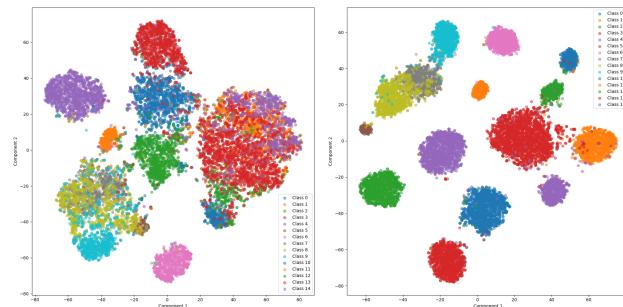
When analyzing the model performance across various datasets, we observe that on the Plant Village dataset, the three architectures achieve high accuracy, precision, and recall, with MobileNetV2 performing slightly better than both GoogLeNet and EfficientNet. This superior performance of all models can likely be attributed to the large training set and the simpler backgrounds in the input images compared to those in other datasets. The less complex backgrounds allow the models to focus more effectively on the key features of the plants, leading to improved performance. Apart from the dataset's reduced background noise, the performance of the MobileNetV2 architecture is also attributed to its use of inverted residuals with linear bottlenecks, which help capture fine-grained features with fewer parameters.

As the complexity of input images increases in the Crop Disease dataset, which features images with varying degrees of background noise across different classes, GoogLeNet tends to outperform both EfficientNet and MobileNetV2 in terms of precision and recall. This superior performance is likely attributable to the inception module within GoogLeNet, which enables the network to capture and integrate multi-scale features more effectively from the same input. This leads to a nuanced understanding of the images, resulting in improved precision and recall metrics. Table 3 illustrates the comparative performance of GoogLeNet and MobileNetV2 when utilizing pre-trained weights on this dataset. Furthermore, as depicted in Figure

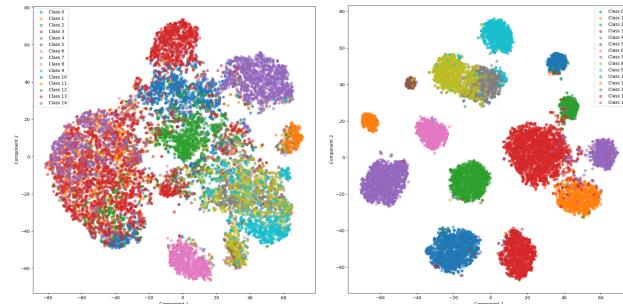
4, MobileNetV2 demonstrates a performance comparable to GoogLeNet, suggesting it as a viable alternative when used with pre-trained weights, especially for deployment on edge devices.

Model	Accuracy (%)	Precision (%)	Recall (%)	F1 Score
GoogleNet	87.90	87.23	88.21	0.88
MobileNet V2	88.68	87.79	88.59	0.88

Table 3. Transfer Learning: Performance of GoogleNet & MobileNet V2 Architectures on Crop Disease dataset



(a) GoogLeNet without Pre-trained Weights (b) GoogLeNet with Pre-trained Weights



(c) MobileNetV2 without Pre-trained Weights (d) MobileNetV2 with Pre-trained Weights

Figure 4. t-SNE for four models on train set of Crop Disease dataset

With input images featuring the highest degree of background noise, as seen in the Cassava dataset, MobileNetV2 outperforms other network architectures, achieving an accuracy of 66.55%, a precision of 58.14%, and a recall of 44.13%. While the high accuracy could be attributed to the utilization of inverted residuals with linear bottlenecks which helps capture fine-grained embedded features in the input, we observe a significant trade-off in precision and recall compared to other architectures. Figure 34 shows the model's confidence in identifying a subset of classes (class 1 and 3) precisely but struggling with others, leading to lower recall. Despite having better performance on the Crop Disease dataset, the relative complexity of GoogLeNet hinders its performance on this smaller dataset, making it challenging to learn nuanced patterns in the images.

486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539



GROUP Q: FINAL REPORT

540

4.3. Ablative Study

In our ablation study, we investigate the impact of employing additional data augmentation techniques, varying the number of images per class, applying transfer learning, and tuning hyperparameters such as batch size and learning rate on the Cassava dataset using the EfficientNet architecture. Among the three datasets examined, we found the Plant Village dataset to be the easiest to classify due to the minimal background noise in the inputs, followed by the Crop Diseases dataset. The Cassava dataset presents the most significant classification challenge, as evidenced by the results presented in Table 2. This difficulty arises from the high level of background noise in the images, which impedes the model's ability to capture essential features. To mitigate the effects of high background noise, we employed center cropping, as seen in Figure 9. Although this transformation reduces background noise by reducing the dimensions of an image while focusing on the central region, which contains the subject of interest, it is not as effective as segmentation and leads to no improvement in EfficientNet's performance.

As we observed an imbalance in the dataset impacting the model's performance, we reduced the number of images for the oversampled classes. Specifically, we sampled only 15% of the images for the "CMD" class and 50% for the "CBSD" class. Upon retraining the EfficientNet for 20 epochs on this modified dataset with a batch size of 32 and a learning rate of 0.001, we noted that the model was unable to generalize well due to the decreased overall dataset size. This necessitated an increase in the number of epochs and tweaking of hyper-parameters leading to a slower convergence to optimal solution.

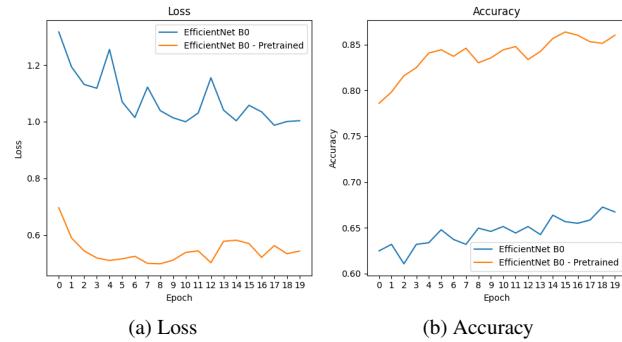


Figure 5. EfficientNet with and without pre-trained weights on Cassava dataset

We also investigated the impact of training EfficientNet with pre-trained weights on the original Cassava dataset through a transfer learning approach. When equipped with pre-trained weights, EfficientNet achieved an accuracy of 84.78%, a precision of 80.91%, and a recall of 78.31%. As depicted in Figure 5, the model benefits from the use of pre-

Batch Size	Learning Rate	Loss	Accuracy (%)
8	0.0013	1.53	62.48
32	0.0095	4.90	53.99
16	0.0001	1.07	61.95
32	0.0039	1.04	66.02
8	0.0086	1.16	59.23
64	0.0026	0.94	65.13
8	0.0096	3.66	46.19
64	0.0095	1.38	61.59
16	0.0028	1.77	53.63
8	0.0021	16.54	52.74

Table 4. Summary of Hyper-parameter Tuning Trials

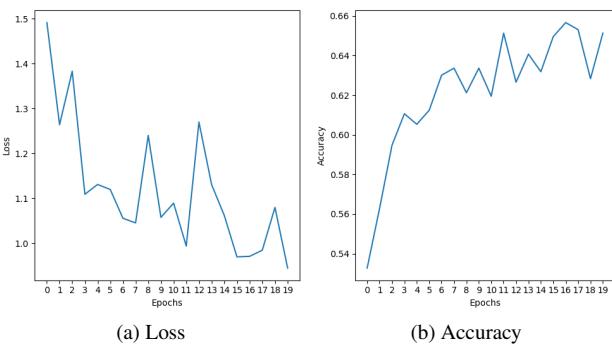


Figure 6. EfficientNet on Cassava dataset with a batch size of 64 and a learning rate of 0.002621

trained weights, which act as powerful feature extractors, resulting in approximately 20% improvement in accuracy over the non-pretrained model. This leads to faster convergence and improved performance across the dataset.

Additionally, we optimized the batch size and learning rate for EfficientNet using the Random Search algorithm. We randomly sampled 10 configurations from the specified batch sizes: 8, 16, 32, 64, with learning rates ranging from 0.0001 to 0.001. The performance of the model under these configurations, as illustrated in Table 4, underscores the critical role of hyperparameter tuning in achieving a balance between training stability and model performance. Our observations indicate that larger batch sizes, particularly 64, when paired with a moderate learning rate of 0.0026, tend to offer an optimal balance by providing training stability and a model loss of less than 1. Conversely, smaller batch sizes combined with moderate to high learning rates (e.g., 0.00028 or 0.0096) appear to induce training instability, resulting in a loss greater than 1. With an optimal configuration of a batch size of 64 and a learning rate of 0.0026, EfficientNet achieved an accuracy of 70.09%, precision of 59.33%, and recall of 46.34% over unseen images of the Cassava dataset.

594

595

596

597

598

599

600

601

602

603

604

605

606

607

608

609

610

611

612

613

614

615

616

617

618

619

620

621

622

623

624

625

626

627

628

629

630

631

632

633

634

635

636

637

638

639

640

641

642

643

644

645

646

647

GROUP Q: FINAL REPORT

648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701

References

- [1] Mwebaze, Ernest, et al. "iCassava 2019 Fine-Grained Visual Categorization Challenge Dataset." Google Research, Artificial Intelligence Lab Makerere University, National Crops Resources Research Institute. Accessed [Access Date], <https://tensorflow.google.cn/datasets/catalog/cassava>. 2 702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
- [2] Mensah Kwabena, Patrick, et al. "Dataset for Crop Pest and Disease Detection." Mendeley Data, vol. 1, 2023, doi:10.17632/bwh3zbpkpv.1. <https://data.mendeley.com/datasets/bwh3zbpkpv/1> 2
- [3] Hughes, D. P., and M. Salathé. "An open access repository of images on plant health to enable the development of mobile disease diagnostics through machine learning and crowdsourcing." CoRR, vol. abs/1511.08060, 2015, <http://arxiv.org/abs/1511.08060>. https://www.tensorflow.org/datasets/catalog/plant_village 2
- [4] Mohanty, Sharada P., David P. Hughes, and Marcel Salathé. "Using deep learning for image-based plant disease detection." Frontiers in plant science 7 (2016): 1419.
- [5] J. Deng, W. Dong, R. Socher, L. -J. Li, Kai Li and Li Fei-Fei, "ImageNet: A large-scale hierarchical image database," 2009 IEEE Conference on Computer Vision and Pattern Recognition, Miami, FL, USA, 2009, pp. 248-255. 3
- [6] Szegedy, C., Liu, W., Jia, Y., Sermanet, P., Reed, S., Anguelov, D., ... and Rabinovich, A. (2015). "Going deeper with convolutions". In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 1-9).
- [7] Tan, Mingxing, and Quoc Le. "Efficientnet: Rethinking model scaling for convolutional neural networks." International conference on machine learning. PMLR, 2019.
- [8] Sandler, Mark, et al. "Mobilenetv2: Inverted residuals and linear bottlenecks". Proceedings of the IEEE conference on computer vision and pattern recognition. 2018. 3
- [9] Paszke, Adam, et al. "Pytorch: An imperative style, high-performance deep learning library". Advances in neural information processing systems 32 (2019).
- [10] Punam Bedi, Pushkar Gole, Plant Disease Detection Using Hybrid Model Based on Convolutional Autoencoder and Convolutional Neural Network, *Artificial Intelligence in Agriculture*, Volume 5, 2021, Pages 90-101, ISSN 2589-7217.
- [11] Hassan, S.M.; Maji, A.K.; Jasiński, M.; Leonowicz, Z.; Jasińska, E. Identification of Plant-Leaf Diseases Using CNN and Transfer-Learning Approach. *Electronics* 2021, 10, 1388.
- [12] Shelar, Nishant; Shinde, Suraj; Sawant, Shubham; Dhumal, Shreyash; Fakir, Kausar. Plant Disease Detection Using CNN. *ITM Web Conf.* 2022, 44, 03049. 2
- [13] Nandi, Rabindra Palash, Aminul Siddique, Nazmul Zilani, Mohammed. (2022). Device-friendly Guava fruit and leaf disease detection using deep learning. 1, 2
- [14] Boulet J, Foucher S, Théau J, St-Charles PL. Convolutional Neural Networks for the Automatic Identification of Plant Diseases. *Front Plant Sci.* 2019 Jul 23;10:941. doi: 10.3389/fpls.2019.00941. PMID: 31396250; PMCID: PMC6664047. 1, 2
- [15] Sharma, Rahul Singh, Amar .., Kavita Jhanjhi, Noor Masud, Mehedi Jaha, Emad Verma, Prof. (2021). Plant Disease Diagnosis and Image Classification Using Deep Learning. *Computers, Materials and Continua.* 71. 2125-2140. 10.32604/cmc.2022.020017. 2
- [16] Khyeh0719."Pytorch EfficientNet Baseline Train AMP Aug". Kaggle, <https://www.kaggle.com/code/khyeh0719/pytorch-efficientnet-baseline-train-amp-aug>. Accessed 10 April 2024.
- [17] Sik-Ho Tsang. "Review: MobileNetV2 — Light Weight Model for Image Classification." Towards Data Science. <https://towardsdatascience.com/review-mobilenetv2-light-weight-model-image-classification-8febb490e61c>. Accessed 10 April 2024. 3
- [18] Sandler, Mark, et al. "Mobilenetv2: Inverted residuals and linear bottlenecks." Proceedings of the IEEE conference on computer vision and pattern recognition. 2018. 3
- [19] Petru Potrimba. "What is EfficientNet?" Roboflow Blog. <https://blog.roboflow.com/what-is-efficientnet/>. Accessed 10 April 2024. 4
- [20] DhanushKumar. "EfficientNet: Scaling Depth, Width, Resolution." Medium. <https://medium.com/@danushidk507/efficientnet-scaling-depth-width-resolution-11e2d4311357>. Accessed 10 April 2024. 4

GROUP Q: FINAL REPORT

- 756 [21] Tan, Mingxing, and Quoc V. Le. "EfficientNet: Re- 810
757 thinking Model Scaling for Convolutional Neural Net- 811
758 works." International Conference on Machine Learn- 812
759 ing, 2019. 4 813
760
- 761 [22] Sandler, Mark, et al. "Mobilenetv2: Inverted residuals 814
762 and linear bottlenecks." Proceedings of the IEEE con- 815
763 ference on computer vision and pattern recognition. 816
764 2018. 3 817
765
- 766 [23] Szegedy, Christian, et al. "Going deeper with convo- 818
767 lutions." Proceedings of the IEEE conference on com- 819
768 puter vision and pattern recognition. 2015. 3 820
769
- 770 [24] Ruder, Sebastian. "An overview of gradient de- 821
771 scent optimization algorithms." arXiv preprint 822
772 arXiv:1609.04747 (2016). 4 823
773
- 774
- 775
- 776
- 777
- 778
- 779
- 780
- 781
- 782
- 783
- 784
- 785
- 786
- 787
- 788
- 789
- 790
- 791
- 792
- 793
- 794
- 795
- 796
- 797
- 798
- 799
- 800
- 801
- 802
- 803
- 804
- 805
- 806
- 807
- 808
- 809

GROUP Q: FINAL REPORT

864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917

5. Supplementary Material

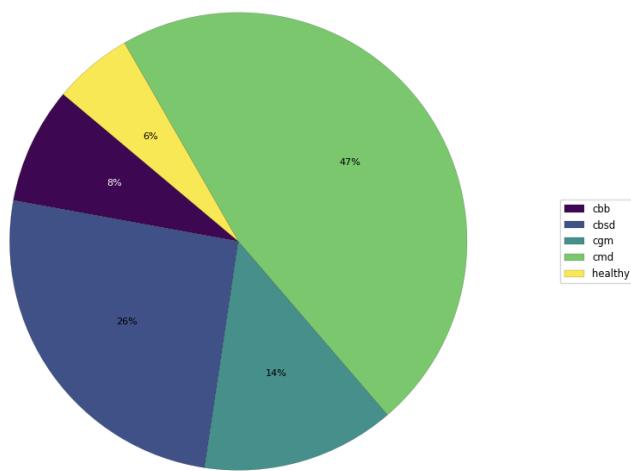


Figure 7. Inter-Class variance of Cassava Dataset

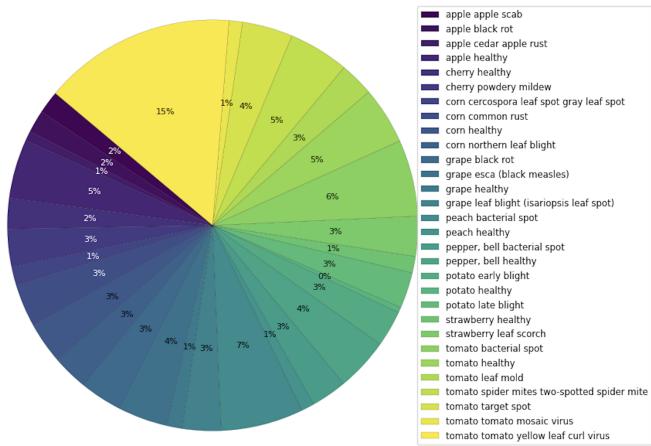


Figure 8. Inter-Class variance of Plant Village Dataset

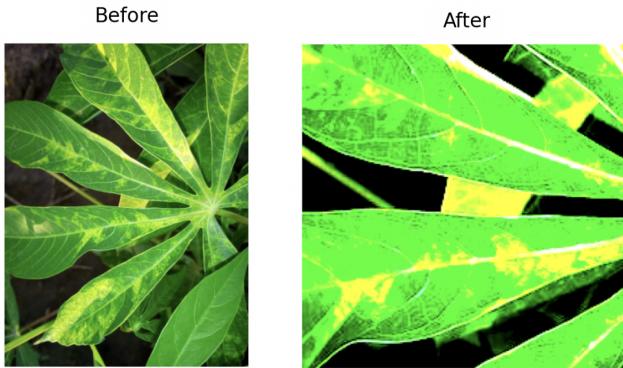


Figure 9. Example of center cropping on Cassava dataset



Figure 10. Samples of Plant Village Dataset



Figure 11. Samples of Crop Disease Dataset

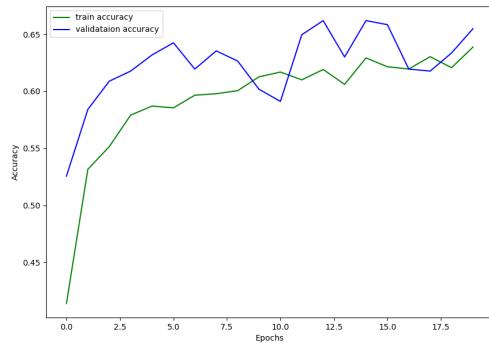


Figure 12. Accuracy of GoogleNet on Cassava Dataset

918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971

GROUP Q: FINAL REPORT

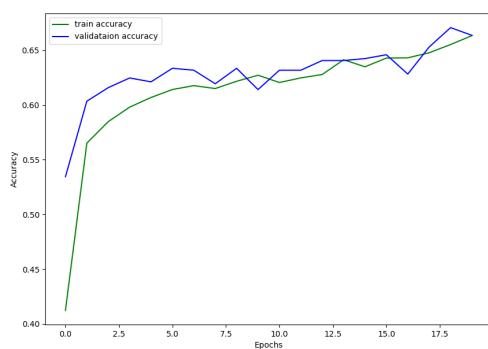


Figure 13. Accuracy of EfficientNet B0 on Cassava Dataset

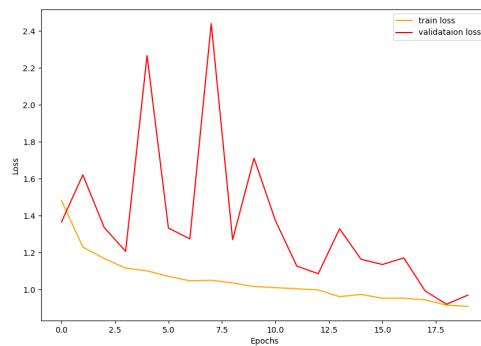


Figure 16. Loss of EfficientNet B0 on Cassava Dataset

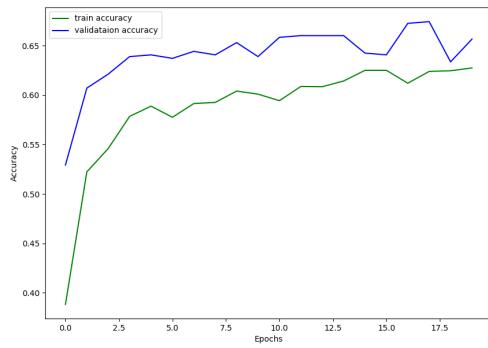


Figure 14. Accuracy of MobileNetV2 on Cassava Dataset

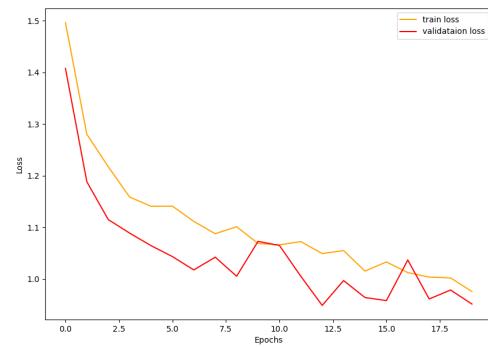


Figure 15. Loss of GoogleNet on Cassava Dataset

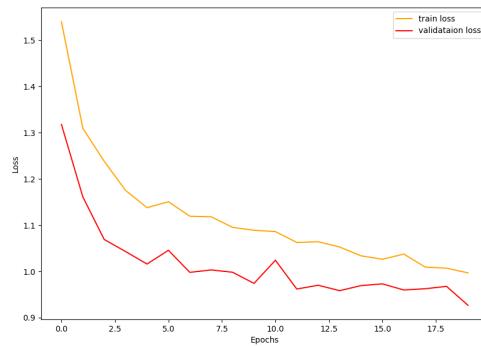


Figure 17. Loss of MobileNetV2 on Cassava Dataset

GROUP Q: FINAL REPORT

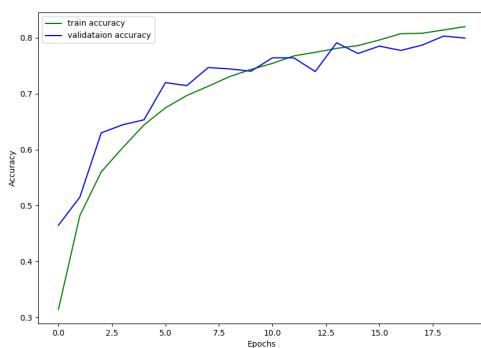


Figure 18. Accuracy of GoogleNet on Crop Disease Dataset

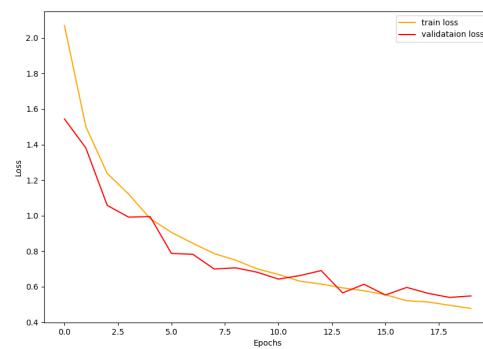


Figure 21. Loss of GoogleNet on Crop Disease Dataset

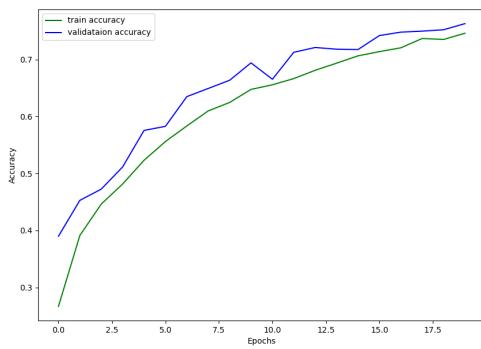


Figure 19. Accuracy of EfficientNet B0 on Crop Disease Dataset

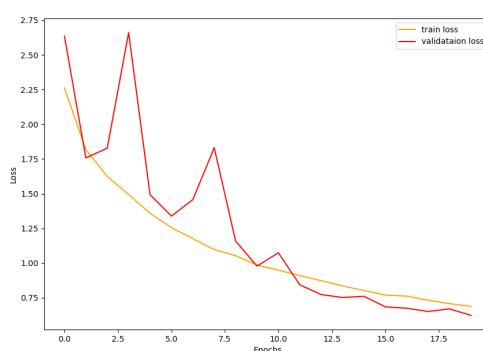


Figure 22. Loss of EfficientNet B0 on Crop Disease Dataset

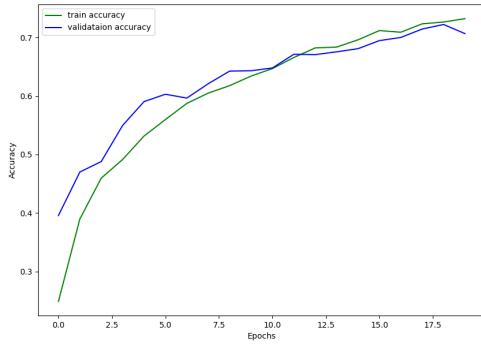


Figure 20. Accuracy of MobileNetV2 on Crop Disease Dataset

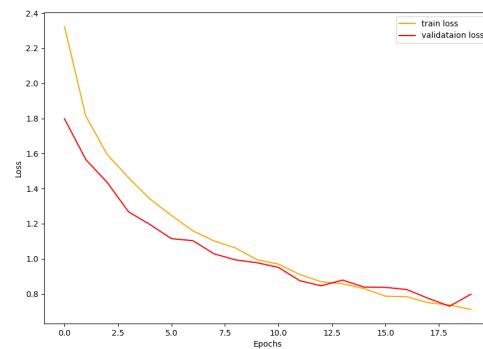


Figure 23. Loss of MobileNetV2 on Crop Disease Dataset

GROUP Q: FINAL REPORT

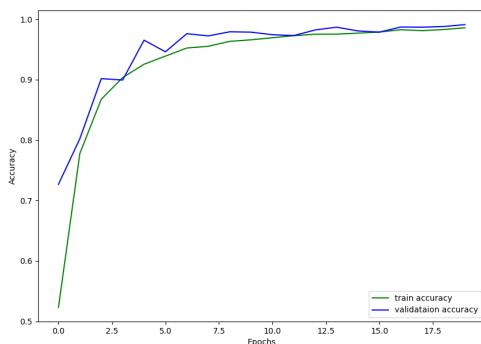


Figure 24. Accuracy of GoogleNet on Plant Village Dataset

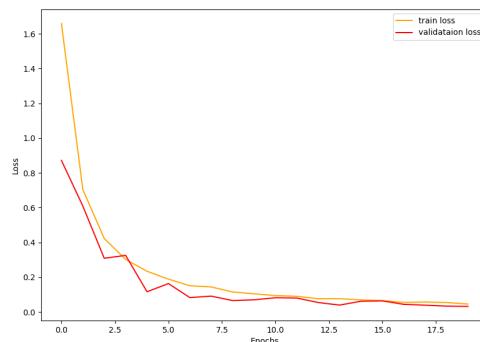


Figure 27. Loss of GoogleNet on Plant Village Dataset

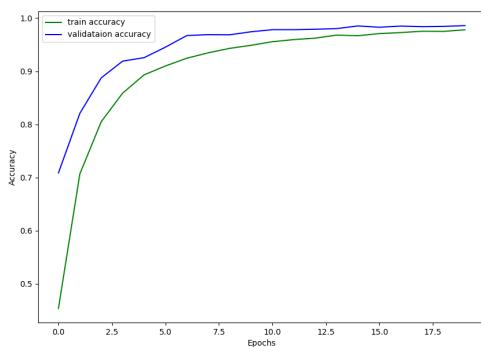


Figure 25. Accuracy of EfficientNet B0 on Plant Village Dataset

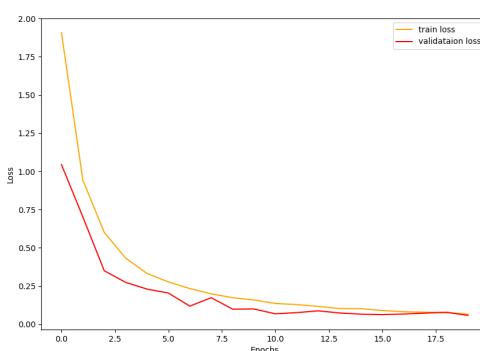


Figure 28. Loss of EfficientNet B0 on Plant Village Dataset

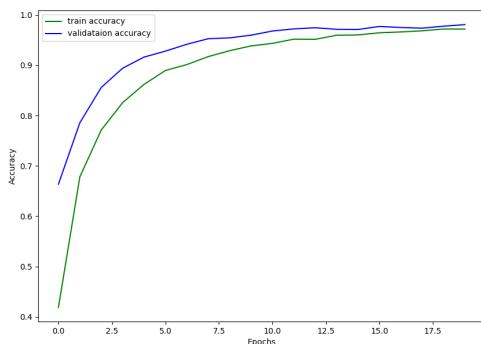


Figure 26. Accuracy of MobileNetV2 on Plant Village Dataset

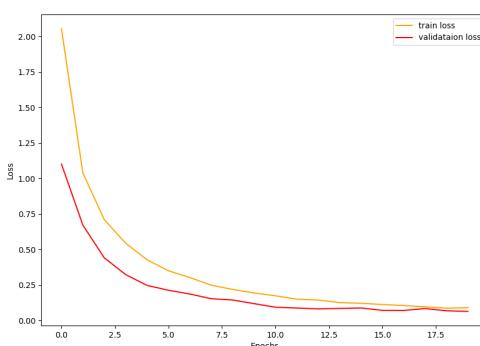
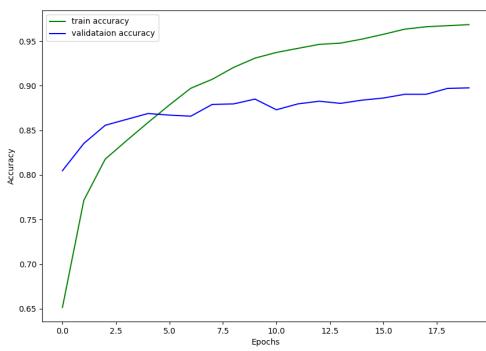


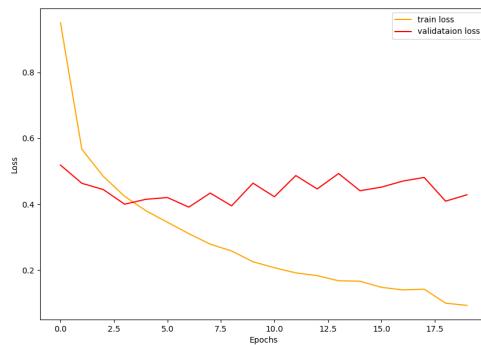
Figure 29. Loss of MobileNetV2 on Plant Village Dataset

1296 **5.1. Transfer Learning**

1350



1312 Figure 30. Accuracy-GoogleNet



1313 Figure 33. Loss-MobileNetV2

1351

1352

1353

1354

1355

1356

1357

1358

1359

1360

1361

1362

1363

1364

1365

1366

1367

1368

1369

1370

1371

1372

1373

1374

1375

1376

1377

1378

1379

1380

1381

1382

1383

1384

1385

1386

1387

1388

1389

1390

1391

1392

1393

1394

1395

1396

1397

1398

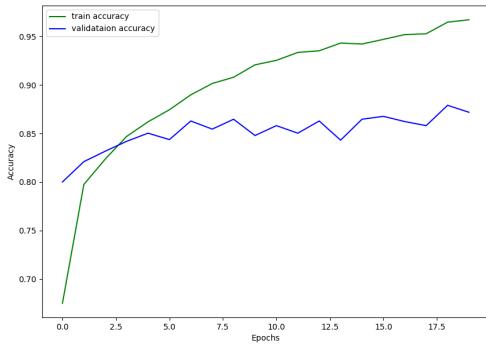
1399

1400

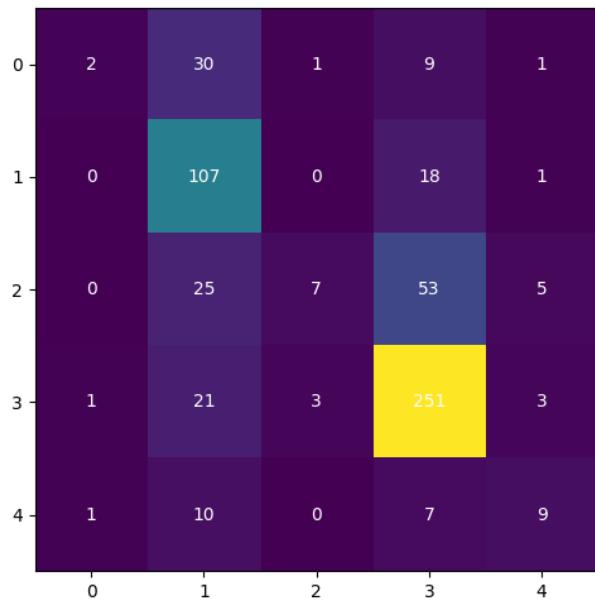
1401

1402

1403



1327 Figure 31. Accuracy-MobileNetV2



1343 Figure 34. Confusion Matrix for MobileNetV2 on Cassava dataset

1397

1398

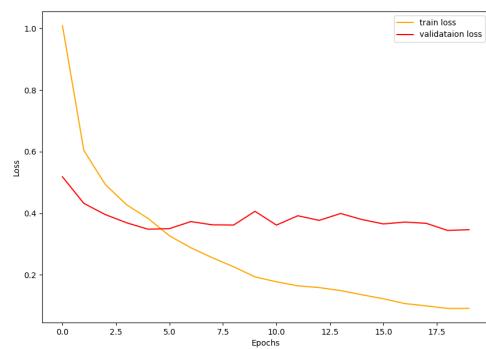
1399

1400

1401

1402

1403



1338 Figure 32. Loss-GoogleNet